

# CS4830 - BIG DATA LABORATORY

## FINAL PROJECT REPORT

**Group Name:** Quarantined Cops

**Members:**

Roll Number	Name
CH16B024	Vishwesh Ramanathan
CH16B033	Ashwin S Nair
CH16B045	Raghav Rajesh Moar

### ➤ Objective:

- 1) To train a machine learning model on big dataset of NYC Parking tickets to predict the violation location
- 2) To be able to perform real-time computation using the trained model using Big data technologies taught in course CS4830.

### ➤ Initial Data Inspection:

No. of Rows/Entries: **1,18,09,233**

No of features (Columns): **50 columns**

16 features have only NaN entries for all the Rows and can be removed without affecting the model

### ➤ Pre-processing:

#### 1) Removing Duplicate Summons numbers:

Since Summons numbers are unique for each row entry and having a duplicate summon implies that full row is repeated. Therefore all repeated rows are removed.

```
data_tickets.iloc[[141927,141928]]
```

	Summons Number	Plate ID	Registration State	Plate Type	Issue Date	Violation Code	Vehicle Body Type	Vehicle Make	Issuing Agency	Street Code1	Street Code2	Street Code3	Vehicle Expiration Date	Violation Count
141927	1370137552	GJY9810	NY	PAS	08/22/2014	21	SDN	TOYOT	S	0	0	0	01/01/20160306 12:00:00 PM	73
141928	1370137552	GJY9810	NY	PAS	08/22/2014	21	SDN	TOYOT	S	0	0	0	01/01/20160306 12:00:00 PM	73

**Fig 1. Duplicate Summon Number implies duplicate Row**

## 2) Removing not important Features:

We inspected every feature manually. We found many features to not correlate well with the violation locations. Some of them are:

**'Summons Number'** as it just indicates new entry and has nothing to with violation location

**'Vehicle Expiration Date'** again its specific to vehicles and nothing indicating about violation location

**'Violation Legal Code'** as it indicates just the type of violation and not specific to some location

Similarly, other 14 not important features are removed. These features are removed based on either having a lot of categorical values with fewer data points for each category or having lots of missing data points. Features having lots of categorical values will only add noise to the modeling step. Hence on this basis, a lot of the features are removed.

**So in total 16+17 = 33 features in total are removed.**

```
#Data preprocessing pipeline
df=df.dropDuplicates(subset=['Summons Number'])
columns_to_drop = ['Plate ID','Issuer Code','Time First Observed','Vehicle Expiration Date','House Number','Street Name','Intersecting Street','[
df = df.drop(*columns_to_drop)
df.show(5)
```

Summons Number	Registration State	Plate Type	Issue Date	Violation Code	Vehicle Body Type	Vehicle Make	Issuing Agency	Street Code1	Street Code2	S
1366539345	NY	PAS	07/01/2014	31	SDN	HONDA	P	34430	10410	
1367471345	NY	PAS	08/09/2014	71	SDN	TOYOT	P	27820	31390	
1368841090	NY	SRF	11/20/2014	40	SUBN	CHEVR	P	8440	20190	
1369109532	NY	COM	06/23/2014	46	VAN	MERCU	P	18770	10010	
1369992300	ME	PAS	08/20/2014	66	TRAI	null	S	62030	81050	

only showing top 5 rows

**Fig 2. Dropping Unnecessary Columns**

(\*will edit image with removing summon number as well and showing all features that are removed)

## 3) Extracting Issue Date (Feature Engineering):

We split the Issue date into Year Month and Day columns. These columns are also converted to integers. We have also added a column, Day of the week which is calculated form year (since it can be an important feature) (using a python function)

The year column is further used to remove years with 2012 or less since this data corresponds to the years 2013-2014 and the data before these years are considered as outliers/ faulty rows.

The Day column is useful as it can help to distinguish between Weekdays and Weekends based on Violation tickets issued at specific locations.

```

#UDFS
#Extracting information from date
def day_finder(x):
    return datetime.datetime.strptime(x, '%m/%d/%Y').weekday()
#Bucketizing violation time
def time_bucket(x):
    #Bucketizing the time into 8 buckets
    if x is None:
        return 3
    if x[-1]!='P' and x[-1]!='A':
        return 3
    try:
        time = int(x[:-1])
    except:
        return 3
    if x[-1]=='P':
        time = 1200+time
    for i in range(8):
        if time>=300*i and time<300*(i+1):
            return i
time_udf = udf(lambda x: time_bucket(x), IntegerType())
day_udf = udf(lambda x: day_finder(x), IntegerType())

```

```

#Splitting the issue date into month,year,day
df_new = df.withColumn('Month',split('Issue Date','/')[0]).withColumn('Year',split('Issue Date','/')[2]).withColumn('Day',day_udf(col('Issue Date'))).withColumn(
#converting the columns into integers
df_new = df_new.withColumn("Year",df_new["Year"].cast(IntegerType())).withColumn("Month",df_new["Month"].cast(DoubleType())).withColumn("Day",df_new["Day"].cast(
#Removing outliers and some filtering
df_new = df_new.where(f.col("Year")>2012)
#Dropping columns
df_new=df_new.drop(*['Issue Date','Violation Time','Year','Issuer Squad'])
#Filling na
df_new = df_new.fillna({'Time':3})
#Removing na localons of violation location and violation count
df_new = df_new.dropna(how='any',subset=['Violation Location','Violation County'])
#Fill na of these columns using respective max values
# cols = ['Vehicle Body Type','Vehicle Make','Violation County','Violation In Front Of Or Opposite']
# agg_expr = [mode(f.collect_list(col)).alias(col) for col in cols]
# max_vals = df_new.agg(*agg_expr).collect()[0]
# df_new = df_new.fillna({'Vehicle Body Type':max_vals['Vehicle Body Type'],'Vehicle Make':max_vals['Vehicle Make'],'Violation County':max_vals['Violation County']
df_new=df_new.dropna(how='any')
#Renaming columns
names = df_new.schema.names
for name in names:
    df_new = df_new.withColumnRenamed(name, name.replace(" ", "_"))
#Mapping violation location
df_new = df_new.withColumn('Violation_Location', regexp_replace('Violation_Location', 'KINGS', 'K'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'KING', 'K'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'QUEEN', 'Q'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'QU', 'Q'))\

```

**Fig 3. Date and Time Feature Engineering**

#### 4) Bucketing time (Feature Engineering):

The Violation time column is put into 8 buckets of 3 hours each as certain time buckets will have more violations and certain localities will have violation tickets issued during specific hours. Later violation time is dropped.

#### 5) Correction in Violation Location

Firstly, we drop all violations Locations with Nan values as this is our target column and we cannot have supervised learning without labels. Some of the target labels have been names differently. Hence, we mapped the differently named labels to the actual prominent label as shown below.

Also, note that we have mapped Manhattan(MH/MAN) to NY because of only 6 training data points being available which won't lead to a good prediction of that class. Since Manhattan also comes under the broad category of New York, we have hence mapped it as proposed.

```
for name in names:
    df_new = df_new.withColumnRenamed(name, name.replace(" ", "_"))
#Mapping violation location
df_new = df_new.withColumn('Violation_Location', regexp_replace('Violation_Location', 'KINGS', 'K'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'KING', 'K'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'QUEEN', 'Q'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'QU', 'Q'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'NEWY', 'NY'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'NEW Y', 'NY'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'MAN', 'NY'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'MH', 'NY'))\
.withColumn('Violation_Location', regexp_replace('Violation_Location', 'BRONX', 'BX'))
```

**Fig 4. Violation Location correction**

#### 6) Label Encoding:

Finally, all string(text) features are encoded using StringIndexer label encoding. For example registration states are Mapped as NY: 0, TX: 7 in StringIndexer.

=====

### ➤ Performance of the best model and Inferences

**We get the best accuracy with the XGBoost Classifier of 99.9% train and test accuracy (See Fig 5).**

(Note: We also tried RandomForestClassifier giving 95.0% test and train accuracy, and logistic regression giving 49% accuracy)

Also, XGBoost trained pretty fast than Random Forest



```
##### XGBOOST Training #####
#Training XGBOOST only on first few data
df_pandas = df_r1.limit(10000).toPandas()
column_names = ['Registration_State_index', 'Plate_Type_index', 'Violation_Code_index', 'Vehicle_Body_Type_index', 'Vehicle_Make_index', 'Issuing_Agency_index', 'Street_Code1_index', 'Street_Code2_index', 'Street_Code3_index',
X_pandas = df_pandas[column_names].values
y_pandas = df_pandas['label'].values
X_train, X_test, y_train, y_test = train_test_split(X_pandas, y_pandas, test_size = 0.2)
xgboost_model = XGBClassifier(max_depth=7, n_estimators=100, objective='multi:softprob')
xgboost_model.fit(X_train, y_train)
#Testing the model
y_pred_train = xgboost_model.predict(X_train)
print("For train\nAccuracy score:{}, Balanced accuracy score:{}, f1 score:{}".format(accuracy_score(y_train, y_pred_train), balanced_accuracy_score(y_train, y_pred_train), f1_score(y_train, y_pred_train, average='weighted')))
y_pred_test = xgboost_model.predict(X_test)
print("For test\nAccuracy score:{}, Balanced accuracy score:{}, f1 score:{}".format(accuracy_score(y_test, y_pred_test), balanced_accuracy_score(y_test, y_pred_test), f1_score(y_test, y_pred_test, average = 'weighted')))
#For saving model
filename = 'XGB_final_model_v1.pkl'
pickle.dump(xgboost_model, open(filename, 'wb'))

For train
Accuracy score:0.9993125, Balanced accuracy score:0.9986929698062774, f1 score:0.9993124867241515
For test
Accuracy score:0.99865, Balanced accuracy score:0.9966405271390689, f1 score:0.9986498110888233
```

**Fig 5. Checking accuracy scores for XGBoost model**

```
#Defining XGBOOST prediction UDF
@f.pandas_udf(returnType=DoubleType())
def predict_pandas_udf(*cols):
    # cols will be a tuple of pandas.Series here.
    X = pd.concat(cols, axis=1)
    return pd.Series(xgboost_model.predict(np.array(X)))
#Get the transformed data
df_full = model.transform(df_new)
#Apply XGBOOST on it
df_r3 = df_full.withColumn('prediction_xgb', predict_pandas_udf(*column_names))
```

**Fig 6: Defining UDF for XGBoost prediction**

```
##### TESTING THE MODEL PERFORMANCE #####
#Train set for random forest
#Create an evaluation object for the model using the R^2 metric
lr_evaluator = MulticlassClassificationEvaluator(predictionCol = "prediction", labelCol="label", metricName="accuracy")
#Print the Evaluation Result
print("Train accuracy for Random forest:{}".format(lr_evaluator.evaluate(df_r1)))

#Test set for random forest
df_test = model.transform(test)
#Create an evaluation object for the model using the R^2 metric
lr_evaluator = MulticlassClassificationEvaluator(predictionCol = "prediction", labelCol="label", metricName="accuracy")
#Print the Evaluation Result
print("Test accuracy for Random forest:{}".format(lr_evaluator.evaluate(df_test)))

#Test set for XGBOOST
#Create an evaluation object for the model using the R^2 metric
lr_evaluator = MulticlassClassificationEvaluator(predictionCol = "prediction_xgb", labelCol="label", metricName="accuracy")
#Print the Evaluation Result
print("Test accuracy for XGBOOST:{}".format(lr_evaluator.evaluate(df_r3)))

Train accuracy for Random forest:0.9454161301650474
Test accuracy for Random forest:0.9439258191842617
Test accuracy for XGBOOST:0.9948276087411125
```

**Fig 7. Accuracy scores of XGBoost and RandomForest**

```
In [7]: #Test set for random forest
df_test = model.transform(test)
#Create an evaluation object for the model using the R^2 metric
lr_evaluator = MulticlassClassificationEvaluator(predictionCol="prediction", labelCol="label", metricName="accuracy")
#Print the Evaluation Result
print("Test accuracy for Random forest:{}\n".format(lr_evaluator.evaluate(df_test)))

Test accuracy for Random forest:0.950346995301
```

Fig 8. Accuracy for Random Forest

## ➤ Real-time computation and latency of processing each window

Even though XGBoost outperformed Random forest by a whopping 4% accuracy increase. However, in real-time, we were facing an issue of loading the saved model of XGBoost from Google Cloud Storage. Hence we have used only Random Forest for real-time computation.

```
import os
from google.cloud import pubsub_v1
import time
from google.cloud import storage

# publisher = pubsub_v1.PublisherClient.from_service_account_json("./gcloudcredentials.json")
publisher = pubsub_v1.PublisherClient()
topic_name = 'projects/big-data-lab-266809/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='to-kafka', # Set this to something appropriate.
)

#publisher.create_topic(topic_name)
# client = storage.Client.from_service_account_json("./gcloudcredentials.json")
client = storage.Client()

bucket = client.get_bucket("ch16b024")
blob = bucket.get_blob("small_temp.csv")
print("Loading data...")
x = blob.download_as_string()
x = x.decode('utf-8')
data = x.split('\n')
print("Done. Pushing data to kafka server...")
for lines in data[1:]:
    if len(lines)==0:
        break
    #Sleeps for 10 seconds
    time.sleep(10)
    publisher.publish(topic_name, lines.encode(), spam=lines)
```

Fig 7. Publish function

```

kafka_topic = 'from-pubsub'
zk = '10.138.0.3:2181'
app_name = 'from-pubsub' # Can be some other name
sc = SparkContext(appName="KafkaPubsub")
ssc = StreamingContext(sc, 30)
sc.setLogLevel("FATAL")

kafkaStream = KafkaUtils.createStream(ssc, zk, app_name, {kafka_topic: 1})

def getSparkSessionInstance(sparkConf):
    if ("sparkSessionSingletonInstance" not in globals()):
        globals()["sparkSessionSingletonInstance"] = SparkSession \
            .builder \
            .config(conf=sparkConf) \
            .getOrCreate()
    return globals()["sparkSessionSingletonInstance"]

```

Fig 8a. For Real-time implementation using Kafka stream

```

# Return pd.Series(xgboost_model.predict(np.array(x)))
udfValueToCategory = udf(valueToCategory, StringType())
udfaccuracy = udf(accuracy_calc, IntegerType())
time_udf = udf(time_bucket, IntegerType())
day_udf = udf(day_finder, IntegerType())
print("##### START #####")

lines = kafkaStream.map(Lambda x: json.loads(x[1])["spam"]).map(Lambda x: x.split(","))
#Loading pipeline
model = PipelineModel.load('gs://ch16b024/model_finalproject_v1/')
#Considered columns
column_names = ['Registration_State_index', 'Plate_Type_index', 'Violation_Code_index',
'Vehicle_Body_Type_index', 'Vehicle_Make_index', 'Issuing_Agency_index', 'Street_Code1_index',
'Street_Code2_index', 'Street_Code3_index', 'Issuer_Precinct_index', 'Issuer_Command_index',
'Violation_In_Front_Of_Or_Opposite_index', 'Violation_County_index', 'Month', 'Day', 'Time']
#Loading xgboost model
# xgboost_model = pickle.load(open("gs://ch16b024/XGB_final_model_v1.pkl", "rb"))
accuracy = 0
completed = 0
def process(rdd):
    start = time.time()
    global accuracy
    global completed
    # Get the singleton instance of SparkSession
    spark = getSparkSessionInstance(rdd.context.getConf())
    # Convert RDD[String] to RDD[Row] to DataFrame
    rowRdd = rdd.map(Lambda x: Row(Summons_Number=str(x[0]), Registration_State=str(x[2]), Plate_Type=str(x[3]),
Violation_Code=str(x[5]), Vehicle_Body_Type=str(x[6]), Vehicle_Make=str(x[7]),
Issuing_Agency=str(x[8]), Street_Code1=str(x[9]), Street_Code2=str(x[10]),
Street_Code3=str(x[11]), Violation_County=str(x[13]),
Issuer_Precinct=str(x[14]), Issuer_Command=str(x[16]), Issuer_Squad=str(x[17]),
Violation_In_Front_Of_Or_Opposite=str(x[21]), Issue_Date=str(x[4]),
Violation_Time=str(x[18]), Violation_Location=str(x[20])))
    df = spark.createDataFrame(rowRdd)

```

Fig 8b. For Real-time implementation using Kafka stream



### job-project29

Start time: Aug 1, 2020, 7:00:11 PM Elapsed time: 21 min 0 sec Status:

Output Configuration

<input type="checkbox"/> Line wrapping												Equivalent command line	
	T101	5	T	PAS	NY	16410	13460	25890	7047383785	SUBN	LEXUS	16	
	0066	66	P	PAS	NY	0	0	0	1372046290	SUBN	HONDA	46	
-----													
Labels correct till now:111/115													
Completed batch of 3 in 4.16602110863sec													
-----													
	Issuer_Command	Issuer_Precinct	Issuing_Agency	Plate_Type	Registration_State	Street_Code1	Street_Code2	Street_Code3	Summons_Number	Vehicle_Body_Type	Vehicle_Make	Violation_Code V	
	T501	1	T	PAS	NY	15350	24750	10510	7818597608	SUBN	NISSA	17	
	T401	108	T	PAS	PA	9140	68290	11390	7008624748	4DSD	VOLVO	40	
	T401	115	T	PAS	NY	16740	8590	8790	7613059390	4DSD	HONDA	38	
-----													
Labels correct till now:113/118													
Completed batch of 3 in 4.7607178688sec													
-----													
	Issuer_Command	Issuer_Precinct	Issuing_Agency	Plate_Type	Registration_State	Street_Code1	Street_Code2	Street_Code3	Summons_Number	Vehicle_Body_Type	Vehicle_Make	Violation_Code V	
	T303	120	T	PAS	NC	44010	49876	17450	7163117346	SUBN	FORD	20	
	T201	52	T	PAS	CT	42820	71220	76790	7773113440	2DSD	HONDA	14	
	0063	63	P	PAS	NY	13980	36830	36880	1381130574	SDN	NISSA	14	
-----													

Fig 9. 118 samples completed after 21 min

### job-project29

Start time: Aug 1, 2020, 7:00:11 PM Elapsed time: 37 min 39 sec Status:

Output Configuration

Line wrapping

	T401	108	T	PAS	NJ	59990	9790	9890	7002227452	TR/C	FORD	38
	T401	108	T	PAS	NJ	59990	9790	9890	7002227452	TR/C	FORD	38

Labels correct till now:205/214

Completed batch of 3 in 4.82632708549sec

	Issuer_Command	Issuer_Precinct	Issuing_Agency	Plate_Type	Registration_State	Street_Code1	Street_Code2	Street_Code3	Summons_Number	Vehicle_Body_Type	Vehicle_Make	Violation_Code V
	T301	73	T	PAS	NY	73470	70930	18830	7134344770	4DSD	NISSA	38
	T302	63	T	PAS	NY	37460	14630	82480	7104778330	4DSD	TOYOT	46
	T401	112	T	PAS	NY	68690	41490	41460	8024520813	4DSD	CHEVR	38

Labels correct till now:208/217

Completed batch of 3 in 4.6467320919sec

	Issuer_Command	Issuer_Precinct	Issuing_Agency	Plate_Type	Registration_State	Street_Code1	Street_Code2	Street_Code3	Summons_Number	Vehicle_Body_Type	Vehicle_Make	Violation_Code V
	T105	5	T	PAS	NJ	32900	24390	14590	7627892790	4DSD	TOYOT	37
	T401	108	T	PAS	NY	10540	10440	10590	8024605788	4DSD	SUBAR	20
	T401	108	T	PAS	NY	10540	10440	10590	8024605788	4DSD	SUBAR	20

Fig 10. 217 samples completed after 38 min

38 min/217samples  $\approx$  0.175 min per sample

21min/118 samples  $\approx$  0.177 min per sample

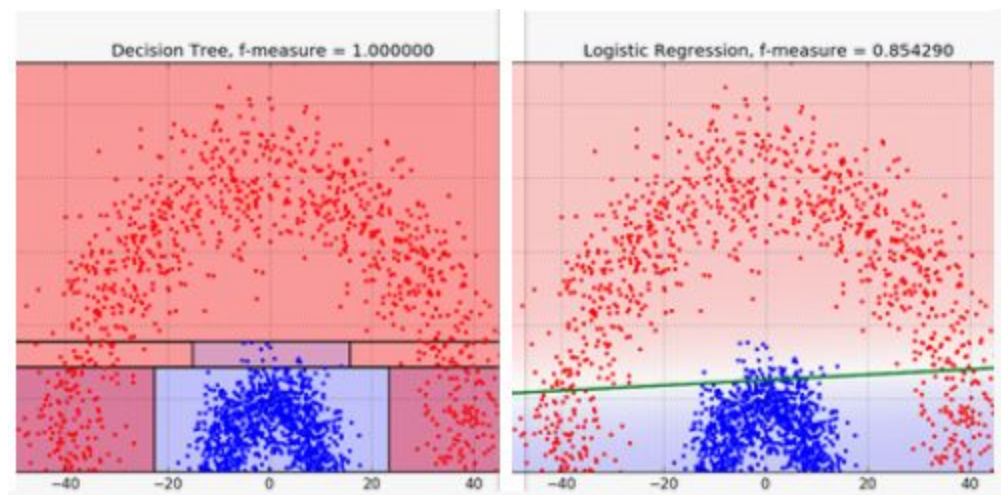


## ➤ Conclusion

- 1) XGBoostClassifier and RandomForestClassifier both perform well (99.9% and 95% accuracy respectively) for multi-classification when feature space is not changed significantly
- 2) XGBoost performed faster than Random Forest
- 3) LogisticRegressionClassifier doesn't perform well for multiclass regression unless we do kernel feature engineering.

### **Reason of why tree-based methods perform well than logistic regression:**

This happens because Tree-based methods bisect the space into several smaller spaces and can, therefore, classify each point correctly given enough tree depth. Whereas logistic regression uses single line space divider and therefore cannot get 100% accuracy even during training if feature space is not transformed



**Fig11. Showing 100% accuracy for tree-based classifier whereas logistic can perform best at F-measure= 0.854**

- 4) Constant latency time for the batch of samples was observed. That is time increases linearly by numbers of samples to evaluate. As streaming of data happens in constant time and processing is fast compared to it.
- 5) Feature Importance graph clearly shows the Violation County Index being the most important feature as each county has a unique index and location(region) that can be guessed correctly based on the county index.

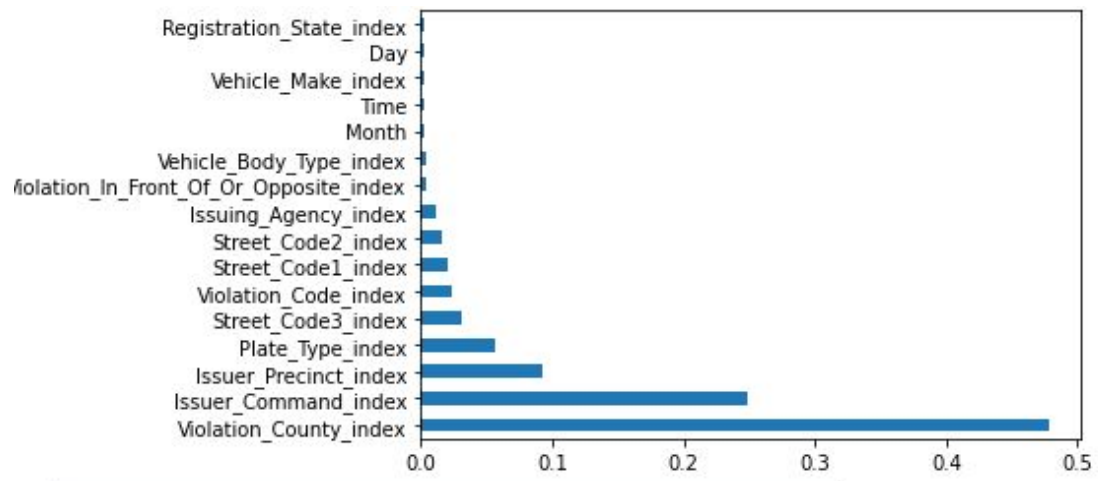


Fig 12. Feature Importance